

**Document Number:** N4071  
**Date:** 2014-06-19  
**Revises:** [N3989](#)  
**Editor:** Jared Hoberock  
NVIDIA Corporation  
[jhoberock@nvidia.com](mailto:jhoberock@nvidia.com)

# **Working Draft, Technical Specification for C++ Extensions for Parallelism**

**Note: this is an early draft. It's known to be incomplet and incorrekt, and it has lots of bad formatting.**

# Contents

<b>1</b>	<b>General</b>	<b>3</b>
1.1	Scope	3
1.2	Normative references	3
1.3	Namespaces and headers	3
1.3.1	Terms and definitions	5
<b>2</b>	<b>Execution policies</b>	<b>6</b>
2.1	In general	6
2.2	Header <experimental/execution_policy> synopsis	6
2.3	Execution policy type trait	7
2.4	Sequential execution policy	7
2.5	Parallel execution policy	8
2.6	Parallel+Vector execution policy	8
2.7	Dynamic execution policy	8
2.7.1	execution_policy construct/assign	9
2.7.2	execution_policy object access	9
2.8	Execution policy objects	10
<b>3</b>	<b>Parallel exceptions</b>	<b>11</b>
3.1	Exception reporting behavior	11
3.2	Header <experimental/exception_list> synopsis	11
<b>4</b>	<b>Parallel algorithms</b>	<b>13</b>
4.1	In general	13
4.1.1	Effect of execution policies on algorithm execution	13
4.1.2	ExecutionPolicy algorithm overloads	14
4.2	Definitions	15
4.3	Non-Numeric Parallel Algorithms	16
4.3.1	Header <experimental/algorithm> synopsis	16
4.3.2	For each	16
4.4	Numeric Parallel Algorithms	17
4.4.1	Header <experimental/numeric> synopsis	17
4.4.2	Reduce	18
4.4.3	Exclusive scan	19
4.4.4	Inclusive scan	20

# 1 General

[parallel.general]

## 1.1 Scope

[parallel.general.scope]

- <sup>1</sup> This Technical Specification describes requirements for implementations of an interface that computer programs written in the C++ programming language may use to invoke algorithms with parallel execution. The algorithms described by this Technical Specification are realizable across a broad class of computer architectures.
- <sup>2</sup> This Technical Specification is non-normative. Some of the functionality described by this Technical Specification may be considered for standardization in a future version of C++, but it is not currently part of any C++ standard. Some of the functionality in this Technical Specification may never be standardized, and other functionality may be standardized in a substantially changed form.
- <sup>3</sup> The goal of this Technical Specification is to build widespread existing practice for parallelism in the C++ standard algorithms library. It gives advice on extensions to those vendors who wish to provide them.

## 1.2 Normative references

[parallel.general.references]

- <sup>1</sup> The following referenced document is indispensable for the application of this document. For dated references, only the edition cited applies. For undated references, the latest edition of the referenced document (including any amendments) applies.  
ISO/IEC 14882:—<sup>1</sup>, *Programming Languages — C++*
- <sup>2</sup> ISO/IEC 14882:— is herein called the *C++ Standard*. The library described in ISO/IEC 14882:— clauses 17-30 is herein called the *C++ Standard Library*. The C++ Standard Library components described in ISO/IEC 14882:— clauses 25 ~~and~~, 26.7 ~~and~~ 20.7.2 are herein called the *C++ Standard Algorithms Library*.
- <sup>3</sup> Unless otherwise specified, the whole of the C++ Standard's Library introduction (C++14 §17) is included into this Technical Specification by reference.

## 1.3 Namespaces and headers

[parallel.general.namespaces]

- <sup>1</sup> Since the the extensions described in this Technical Specification are experimental and not part of the C++ Standard Library, they should not be declared directly within namespace `std`. Unless otherwise specified, all components described in this Technical Specification are declared in namespace `std::experimental::parallel_v1`.  
[ *Note*: Once standardized, the components described by this Technical Specification are expected to be promoted to namespace `std`. — *end note* ]
- <sup>2</sup> Unless otherwise specified, references to such entities described in this Technical Specification are assumed to be qualified with `std::experimental::parallel_v1`, and references to entities described in the C++ Standard Library are assumed to be qualified with `std::`.
- <sup>3</sup> Extensions that are expected to eventually be added to an existing header `<meow>` are provided inside the `<experimental/meow>` header, which shall include the standard contents of `<meow>` as if by  

```
#include <meow>
```

1. To be published. Section references are relative to [N3797](#).

### 1.3.1 Terms and definitions

[parallel.general.defns]

- <sup>1</sup> For the purposes of this document, the terms and definitions given in the C++ Standard and the following apply.
- <sup>2</sup> A *parallel algorithm* is a function template described by this Technical Specification declared in namespace `std::experimental::parallel_v1` with a formal template parameter named `ExecutionPolicy`.

## 2 Execution policies

**[parallel.execpol]**

### 2.1 In general

**[parallel.execpol.general]**

- <sup>1</sup> This clause describes classes that represent *execution policies*. An *execution policy* is an object that expresses the requirements on the ordering of functions invoked as a consequence of the invocation of a standard algorithm. Execution policies afford standard algorithms the discretion to execute in parallel.

[ *Example*:

```
std::vector<int> v = ...

// standard sequential sort
std::sort(vec.begin(), vec.end());
std::sort(std::begin(vec), std::end(vec));
std::sort(vec.begin(), vec.end());

using namespace std::experimental::parallel_v1;

// explicitly sequential sort
sort(seq, v.begin(), v.end());
sort(seq, std::begin(v), std::end(v));
sort(seq, v.begin(), v.end());

// permitting parallel execution
sort(par, v.begin(), v.end());
sort(par, std::begin(v), std::end(v));
sort(par, v.begin(), v.end());

// permitting vectorization as well
sort(vec, v.begin(), v.end());
sort(vec, std::begin(v), std::end(v));
sort(par_vec, v.begin(), v.end());

// sort with dynamically-selected execution
size_t threshold = ...
execution_policy exec = seq;
if (v.size() > threshold)
{
    exec = par;
}

sort(exec, v.begin(), v.end());
sort(exec, std::begin(v), std::end(v));
sort(exec, v.begin(), v.end());
```

— *end example* ]

[ *Note*: Because different parallel architectures may require idiosyncratic parameters for efficient execution, implementations of the Standard Library **should** provide additional execution policies to those described in this Technical Specification as extensions. — *end note* ]

### 2.2 Header <experimental/execution\_policy> synopsis

**[parallel.execpol.synopsis]**

```
namespace std {
namespace experimental {
inline namespace parallel_v1 {
    // 2.3, Execution policy type trait
    template<class T> struct is_execution_policy;
```

```

template<class T> constexpr bool is_execution_policy_v = is_execution_policy<T>::value;

// 2.4, Sequential execution policy
class sequential_execution_policy;

// 2.5, Parallel execution policy
class parallel_execution_policy;

// 2.6, Parallel+Vector execution policy
class vector_execution_policyparallel_vector_execution_policy;

// 2.7, Dynamic execution policy
class execution_policy;
}
}
}

```

## 2.3 Execution policy type trait

[parallel.execpol.type]

```

namespace std {
namespace experimental {
namespace parallel_v1 {
    template<class T> struct is_execution_policy { { see below };
        :: integral_constant<bool, see below> { };
};
};
}

```

- <sup>1</sup> is\_execution\_policy can be used to detect parallel execution policies for the purpose of excluding function signatures from otherwise ambiguous overload resolution participation.

~~If T is the type of a standard or implementation-defined execution policy, is\_execution\_policy<T> shall be publicly derived from integral\_constant<bool, true>, otherwise from integral\_constant<bool, false>.~~

~~is\_execution\_policy<T> shall be a UnaryTypeTrait with a BaseCharacteristic of true\_type if T is the type of a standard or implementation-defined execution policy, otherwise false\_type.~~

~~The behavior of a program that adds specializations for is\_execution\_policy is undefined.~~

## 2.4 Sequential execution policy

[parallel.execpol.seq]

```

namespace std {
namespace experimental {
namespace parallel_v1 {
    class sequential_execution_policy{ unspecified };
};
};
}

```

- <sup>1</sup> The class sequential\_execution\_policy is an execution policy type used as a unique type to disambiguate parallel algorithm overloading and require that a parallel algorithm's execution may not be parallelized.

## 2.5 Parallel execution policy

[\[parallel.execpol.par\]](#)

```
namespace std {
namespace experimental {
namespace parallel_v1 {

    class parallel_execution_policy{ unspecified };

}
}
}
```

- <sup>1</sup> The class `parallel_execution_policy` is an execution policy type used as a unique type to disambiguate parallel algorithm overloading and indicate that a parallel algorithm's execution may be parallelized.

## 2.6 **Parallel+Vector** execution policy

[\[parallel.execpol.vec\]](#)

```
namespace std {
namespace experimental {
namespace parallel_v1 {

    class vector_execution_policyparallel_vector_execution_policy{ unspecified };

}
}
}
```

- <sup>1</sup> The class `class vector_execution_policyparallel_vector_execution_policy` is an execution policy type used as a unique type to disambiguate parallel algorithm overloading and indicate that a parallel algorithm's execution may be vectorized and parallelized.

## 2.7 Dynamic execution policy

[\[parallel.execpol.dynamic\]](#)

```
namespace std {
namespace experimental {
namespace parallel_v1 {

    class execution_policy
    {
    public:
        // 2.7.1, execution_policy construct/assign
        template<class T> execution_policy(const T& exec);
        template<class T> execution_policy& operator=(const T& exec);

        // 2.7.2, execution_policy object access
        template<class T> T* get() noexcept;
        template<class T> const T* get() const noexcept;
    };

}
}
}
```

- <sup>1</sup> The class `execution_policy` is a container for execution policy objects. `execution_policy` allows dynamic control over standard algorithm execution.

[ *Example:*



```

std::vector<float> sort_me = ...

using namespace std::experimental::parallel;
std::execution_policy exec = std::seq;

if(sort_me.size() > threshold)
{
    exec = std::par;
}

std::sort(exec, sort_me.begin(), sort_me.end());
std::sort(exec, std::begin(sort_me), std::end(sort_me));

```

— *end example* ]

- <sup>2</sup> Objects of type `execution_policy` shall be constructible and assignable from objects of type `T` for which `is_execution_policy<T>::value` is true.

### 2.7.1 `execution_policy` construct/assign

[\[parallel.execpol.con\]](#)

- 1 `template<class T> execution_policy(const T& exec);`

<sup>2</sup> *Effects:* Constructs an `execution_policy` object with a copy of `exec`'s state.

*Requires:* `is_execution_policy<T>::value` is true.

*Remarks:* This constructor shall not participate in overload resolution unless `is_execution_policy<T>::value` is true.

- 3 `template<class T> execution_policy& operator=(const T& exec);`

<sup>4</sup> *Effects:* Assigns a copy of `exec`'s state to `*this`.

<sup>5</sup> *Returns:* `*this`.

*Requires:* `is_execution_policy<T>::value` is true.

*Remarks:* This operator shall not participate in overload resolution unless `is_execution_policy<T>::value` is true.

### 2.7.2 `execution_policy` object access

[\[parallel.execpol.access\]](#)

- 1 `const type_info& type() const noexcept;`

<sup>2</sup> *Returns:* `typeid(T)`, such that `T` is the type of the execution policy object contained by `*this`.

- 3 `template<class T> T* get() noexcept;`  
`template<class T> const T* get() const noexcept;`

<sup>4</sup> *Returns:* If `target_type() == typeid(T)`, a pointer to the stored execution policy object; otherwise a null pointer.

*Requires:* `is_execution_policy<T>::value` is true.

*Remarks:* This function shall not participate in overload resolution unless `is_execution_policy<T>` is true.

## 2.8 Execution policy objects

**[parallel.execpol.objects]**

```

namespace std {
namespace experimental {
namespace parallel_v1 {

    —constexpr sequential_execution_policy seq = sequential_execution_policy();
    —constexpr parallel_execution_policy par = parallel_execution_policy();
    —constexpr vector_execution_policy vec = vector_execution_policy();

    —constexpr sequential_execution_policy seq{};
    —constexpr parallel_execution_policy par{};
    —constexpr parallel_vector_execution_policy par_vec{};

}
}
}

```

- <sup>1</sup> The header `<experimental/execution_policy>` declares a global object associated with each type of execution policy defined by this Technical Specification.

## 3 Parallel exceptions

[\[parallel.exceptions\]](#)

### 3.1 Exception reporting behavior

[\[parallel.exceptions.behavior\]](#)

- <sup>1</sup> ~~During the execution of a standard parallel algorithm, if~~ temporary memory resources are required ~~by the algorithm~~ and none are available, the algorithm throws a `std::bad_alloc` exception.
- <sup>2</sup> During the execution of a standard parallel algorithm, if the application of a function object terminates with an uncaught exception, the behavior of the program is determined by the type of execution policy used to invoke the algorithm:

If the execution policy object is of type class

~~vector\_execution\_policy~~`parallel_vector_execution_policy`, `std::terminate` shall be called.

If the execution policy object is of type `sequential_execution_policy` or

`parallel_execution_policy`, the execution of the algorithm terminates with an

`exception_list` exception. All uncaught exceptions thrown during the application of user-provided function objects shall be contained in the `exception_list`.

[ *Note:* For example, the number of invocations of the user-provided function object in `for_each` is unspecified. When `for_each` is executed sequentially, only one exception will be contained in the `exception_list` object. — *end note* ]

[ *Note:* These guarantees imply that, unless the algorithm has failed to allocate memory and terminated with `std::bad_alloc`, all exceptions thrown during the execution of the algorithm are communicated to the caller. It is unspecified whether an algorithm implementation will "forge ahead" after encountering and capturing a user exception. — *end note* ]

[ *Note:* The algorithm may terminate with the `std::bad_alloc` exception even if one or more user-provided function objects have terminated with an exception. For example, this can happen when an algorithm fails to allocate memory while creating or adding elements to the `exception_list` object. — *end note* ]

If the execution policy object is of any other type, the behavior is implementation-defined.

<sup>3</sup>

### 3.2 Header `<experimental/exception_list>` synopsis

[\[parallel.exceptions.synopsis\]](#)

```
namespace std {
namespace experimental {
inline namespace parallel_v1 {

    class exception_list : public exception
    {
    public:

        typedef exception_ptr                                value_type;
        typedef const value_type&                             reference;
        typedef const value_type&                             const_reference;
        typedef implementation_defined                       const_iterator;
        typedef const_iterator                                iterator;
        typedef typename iterator_traits<const_iterator>::difference_type difference_type;
        typedef size_t                                        size_type;

        typedef unspecified                                   iterator;
    }
```

```
size_t size() const noexcept;
iterator begin() const noexcept;
iterator end() const noexcept;
```

```
const char* what() const override noexcept;
```

```
private:
std::list<exception_ptr> exceptions_; // exposition only
};
}
}
}
```

<sup>1</sup> The class `exception_list` ~~is a container~~ owns a sequence of `exception_ptr` objects. ~~The~~ parallel algorithms may use the `exception_list` to communicate uncaught exceptions encountered during parallel execution to the caller of the algorithm.

<sup>2</sup> The type `exception_list::iterator` shall fulfill the requirements of `ForwardIterator`.

<sup>3</sup> `size_t size() const noexcept;`

<sup>4</sup> *Returns:* The number of `exception_ptr` objects contained within the `exception_list`.

<sup>5</sup> *Complexity:* Constant time.

<sup>6</sup> `exception_list::iterator begin() const noexcept;`

<sup>7</sup> *Returns:* An iterator referring to the first `exception_ptr` object contained within the `exception_list`.

<sup>8</sup> `exception_list::iterator end() const noexcept;`

<sup>9</sup> *Returns:* An iterator ~~which is the past-the-end value for the `exception_list`~~ that is past the end of the owned sequence.

<sup>10</sup> `const char* what() const noexcept;`

<sup>11</sup> *Returns:* An implementation-defined NTBS.

## 4 Parallel algorithms

[\[parallel.alg\]](#)

### 4.1 In general

[\[parallel.alg.general\]](#)

This clause describes components that C++ programs may use to perform operations on containers and other sequences in parallel.

#### 4.1.1 Effect of execution policies on algorithm execution

[\[parallel.alg.general.exec\]](#)

- <sup>1</sup> Parallel algorithms have template parameters named `ExecutionPolicy` which describe the manner in which the execution of these algorithms may be parallelized and the manner in which they apply user-provided function objects.
- <sup>2</sup> The applications of function objects in parallel algorithms invoked with an execution policy object of type `sequential_execution_policy` execute in sequential order in the calling thread.
- <sup>3</sup> The applications of function objects in parallel algorithms invoked with an execution policy object of type `parallel_execution_policy` are permitted to execute in an unordered fashion in unspecified threads, and indeterminately sequenced within each thread. [ *Note:* It is the caller's responsibility to ensure correctness, for example that the invocation does not introduce data races or deadlocks. — *end note* ]

[ *Example:*

```
using namespace std::experimental::parallel;
int a[] = {0,1};
std::vector<int> v;
for_each(par, std::begin(a), std::end(a), [&](int i) {
    v.push_back(i*2+1);
});
foo-bar
```

The program above has a data race because of the unsynchronized access to the container `v`.

— *end example* ]

[ *Example:*

```
using namespace std::experimental::parallel;
std::atomic<int> x = 0;
int a[] = {1,2};
for_each(par, std::begin(a), std::end(a), [&](int n) {
    x.fetch_add(1, std::memory_order_relaxed);
    // spin wait for another iteration to change the value of x
    while (x.load(std::memory_order_relaxed) == 1) { }
});
```

The above example depends on the order of execution of the iterations, and is therefore undefined (may deadlock). — *end example* ]

[ *Example:*

```
using namespace std::experimental::parallel;
int x=0;
std::mutex m;
int a[] = {1,2};
for_each(par, std::begin(a), std::end(a), [&](int) {
    m.lock();
    ++x;
    m.unlock();
});
```

The above example synchronizes access to object `x` ensuring that it is incremented correctly.  
 — *end example* ]

- 4 The applications of function objects in parallel algorithms invoked with an execution policy of type `vector_execution_policyparallel_vector_execution_policy` are permitted to execute in an unordered fashion in unspecified threads, and unsequenced within each thread. [ *Note: This means that multiple function object invocations may be interleaved on a single thread.* — *end note* ] [ *Note: As a consequence, function objects governed by the `vector_execution_policyparallel_vector_execution_policy` policy must not synchronize with each other. Specifically, they must not acquire locks.* — *end note* ]

[ *Example:*

```
using namespace std::experimental::parallel_v1;
int x=0;
std::mutex m;
int a[] = {1,2};
for_each(par_vec, std::begin(a), std::end(a), [&](int) {
    m.lock();
    ++x;
    m.unlock();
});
```

The above program is invalid because the applications of the function object are not guaranteed to run on different threads. — *end example* ]

[ *Note: The application of the function object may result in two consecutive calls to `m.lock` on the same thread, which may deadlock.* — *end note* ]

[ *Note: The semantics of the `parallel_execution_policy` or the `vector_execution_policyparallel_vector_execution_policy` invocation allow the implementation to fall back to sequential execution if the system cannot parallelize an algorithm invocation due to lack of resources.* — *end note* ]

A parallel algorithm invoked with an execution policy object of type `parallel_execution_policy` or `vector_execution_policyparallel_vector_execution_policy` may apply iterator member functions of a stronger category than its specification requires, if such iterators exist. In this case, the application of these member functions are subject to provisions 3. and 4. above, respectively.

[ *Note: For example, an algorithm whose specification requires `InputIterator` but receives a concrete iterator of the category `RandomAccessIterator` may use `operator[]`. In this case, it is the algorithm caller's responsibility to ensure `operator[]` is race-free.* — *end note* ]

- 5 Algorithms invoked with an execution policy object of type `execution_policy` execute internally as if invoked with instances of type `sequential_execution_policy`, `parallel_execution_policy`, or an implementation-defined execution policy type depending on the dynamic value of the `execution_policy` object. the contained execution policy object.
- 6 The semantics of parallel algorithms invoked with an execution policy object of implementation-defined type are `unspecifiedimplementation-defined`.

#### 4.1.2 ExecutionPolicy algorithm overloads

[[parallel.alg.overloads](#)]

- 1 ~~Parallel algorithms coexist alongside their sequential counterparts as overloads distinguished by a formal template parameter named `ExecutionPolicy`. This template parameter corresponds to the parallel algorithm's first function parameter, whose type is `ExecutionPolicy` is the first template parameter and corresponds to the parallel algorithm's first function parameter, whose type is `ExecutionPolicy&&`. The Parallel Algorithms Library provides overloads for each of the algorithms named in Table 1, corresponding to the algorithms with the same name in the C++ Standard Algorithms Library. For each algorithm in Table 1, if there are overloads for~~

corresponding algorithms with the same name in the C++ Standard Algorithms Library, the overloads shall have an additional template type parameter named `ExecutionPolicy&&`, which shall be the first template parameter.

- 2 Unless otherwise specified, the semantics of `ExecutionPolicy` algorithm overloads are identical to their overloads without.
- 3 Parallel algorithms ~~have the requirement~~ `is_execution_policy<ExecutionPolicy>::value is true` shall not participate in overload resolution unless `is_execution_policy<decay_t<ExecutionPolicy>>::value is true`.

The algorithms listed in Table 1 shall have `ExecutionPolicy` overloads.

Table 1 — Table of parallel algorithms

<u>adjacent_difference</u>	<code>adjacent_find</code>	<code>all_of</code>	<code>any_of</code>
<code>copy</code>	<code>copy_if</code>	<code>copy_n</code>	<code>count</code>
<code>count_if</code>	<code>equal</code>	<code>exclusive_scan</code>	<code>fill</code>
<code>fill_n</code>	<code>find</code>	<code>find_end</code>	<code>find_first_of</code>
<code>find_if</code>	<code>find_if_not</code>	<code>for_each</code>	<code>for_each_n</code>
<code>generate</code>	<code>generate_n</code>	<code>includes</code>	<code>inclusive_scan</code>
<u>inner_product</u>	<code>inplace_merge</code>	<code>is_heap</code>	<code>is_heap_until</code>
<code>is_partitioned</code>	<code>is_sorted</code>	<code>is_sorted_until</code>	<code>lexicographical_compare</code>
<code>max_element</code>	<code>merge</code>	<code>min_element</code>	<code>minmax_element</code>
<code>mismatch</code>	<code>move</code>	<code>none_of</code>	<code>nth_element</code>
<code>partial_sort</code>	<code>partial_sort_copy</code>	<code>partition</code>	<code>partition_copy</code>
<code>reduce</code>	<code>remove</code>	<code>remove_copy</code>	<code>remove_copy_if</code>
<code>remove_if</code>	<code>replace</code>	<code>replace_copy</code>	<code>replace_copy_if</code>
<code>replace_if</code>	<code>reverse</code>	<code>reverse_copy</code>	<code>rotate</code>
<code>rotate_copy</code>	<code>search</code>	<code>search_n</code>	<code>set_difference</code>
<code>set_intersection</code>	<code>set_symmetric_difference</code>	<code>set_union</code>	<code>sort</code>
<code>stable_partition</code>	<code>stable_sort</code>	<code>swap_ranges</code>	<code>transform</code>
<code>uninitialized_copy</code>	<code>uninitialized_copy_n</code>	<code>uninitialized_fill</code>	<code>uninitialized_fill_n</code>
<code>unique</code>	<code>unique_copy</code>		

[ Note: Not all algorithms in the Standard Library have counterparts in Table 1. — end note ]

## 4.2 Definitions

[parallel.alg.defns]

- 1 Define `GENERALIZED_SUM`(`op`, `a1`, ..., `aN`) as follows:

a1-when `N` is 1  
`op`(`GENERALIZED_SUM`(`op`, `b1`, ..., `bM`), `GENERALIZED_SUM`(`op`, `bM`, ..., `bN`)) where  
`b1`, ..., `bN` may be any permutation of `a1`, ..., `aN` and  
~~`0 <= M <= N`~~  
~~`1 <= K+1 = M <= N`~~

- 2 Define `GENERALIZED_NONCOMMUTATIVE_SUM`(`op`, `a1`, ..., `aN`) as follows:

a1-when `N` is 1  
`op`(`GENERALIZED_NONCOMMUTATIVE_SUM`(`op`, `a1`, ..., `aM`), `GENERALIZED_NONCOMMUTATIVE_SUM`(`op`, `aM`, ..., `aN`)) where ~~`0 <= M <= N`~~ ~~`1 <= K+1 = M <= N`~~

4

## 4.3 Non-Numeric Parallel Algorithms

[parallel.alg.ops]

### 4.3.1 Header <experimental/algorithm> synopsis

[parallel.alg.ops.synopsis]

```
namespace std {
namespace experimental {
  inline namespace parallel_v1 {
    template<class ExecutionPolicy,
             class InputIterator, class Function>
    void for_each(ExecutionPolicy&& exec,
                  InputIterator first, InputIterator last,
                  Function f);
    template<class InputIterator, class Size, class Function>
    InputIterator for_each_n(InputIterator first, Size n,
                             Function f);
  }
}
}
```

### 4.3.2 For each

[parallel.alg.foreach]

```
1      template<class ExecutionPolicy,
              class InputIterator, class Function>
      void for_each(ExecutionPolicy&& exec,
                    InputIterator first, InputIterator last,
                    Function f);
```

2 *Effects:* Applies  $f$  to the result of dereferencing every iterator in the range  $[first, last)$ .  
 [ *Note:* If the type of  $first$  satisfies the requirements of a mutable iterator,  $f$  may apply nonconstant functions through the dereferenced iterator. — *end note* ]

3 *Complexity:* Applies  $f$  exactly  $last - first$  times.

4 *Remarks:* If  $f$  returns a result, the result is ignored.

5 *Notes:* Unlike its sequential form, the parallel overload of `for_each` does not return a copy of its `Function` parameter, since parallelization may not permit efficient state accumulation.

6 *Requires:* Unlike its sequential form, the parallel overload of `for_each` requires `Function` to meet the requirements of `CopyConstructible`, but not `MoveConstructible`.

```
7      template<class InputIterator, class Size, class Function>
      InputIterator for_each_n(InputIterator first, Size n,
                              Function f);
```

8 *Requires:* `Function` shall meet the requirements of `MoveConstructible` [ *Note:* `Function` need not meet the requirements of `CopyConstructible`. — *end note* ]

9 *Effects:* Applies  $f$  to the result of dereferencing every iterator in the range  $[first, first + n)$ , starting from  $first$  and proceeding to  $first + n - 1$ . [ *Note:* If the type of  $first$  satisfies the requirements of a mutable iterator,  $f$  may apply nonconstant functions through the dereferenced iterator. — *end note* ]

10 *Returns:*  $first + n$  for non-negative values of  $n$  and  $first$  for negative values.

11 *Remarks:* If  $f$  returns a result, the result is ignored.



```
template<class ExecutionPolicy,
        class InputIterator, class Size, class Function>
InputIterator for_each_n(ExecutionPolicy && exec,
                        InputIterator first, Size n,
                        Function f);
```

*Effects:* Applies *f* to the result of dereferencing every iterator in the range  $[first, first + n)$ , starting from *first* and proceeding to  $first + n - 1$ . [ *Note:* If the type of *first* satisfies the requirements of a mutable iterator, *f* may apply nonconstant functions through the dereferenced iterator. — *end note* ]

*Returns:* *first* + *n* for non-negative values of *n* and *first* for negative values.

*Remarks:* If *f* returns a result, the result is ignored.

*Notes:* Unlike its sequential form, the parallel overload of *for\_each\_n* requires *Function* to meet the requirements of *CopyConstructible*, ~~but not *MoveConstructible*.~~

## 4.4 Numeric Parallel Algorithms

[parallel.alg.numeric]

### 4.4.1 Header <experimental/numeric> synopsis

[parallel.alg.numeric.synopsis]

```
namespace std {
namespace experimental {
inline namespace parallel_v1 {
    template<class InputIterator>
        typename iterator_traits<InputIterator>::value_type
            reduce(InputIterator first, InputIterator last);
    template<class InputIterator, class T>
        T reduce(InputIterator first, InputIterator last, T init);
    template<class InputIterator, class T, class BinaryOperation>
        T reduce(InputIterator first, InputIterator last, T init,
                 BinaryOperation binary_op);

template<class InputIterator, class OutputIterator>
OutputIterator
exclusive_scan(InputIterator first, InputIterator last,
OutputIterator result);

    template<class InputIterator, class OutputIterator,
              class T>
        OutputIterator
            exclusive_scan(InputIterator first, InputIterator last,
                          OutputIterator result,
                          T init);
    template<class InputIterator, class OutputIterator,
              class T, class BinaryOperation>
        OutputIterator
            exclusive_scan(InputIterator first, InputIterator last,
                          OutputIterator result,
                          T init, BinaryOperation binary_op);

    template<class InputIterator, class OutputIterator>
        OutputIterator
            inclusive_scan(InputIterator first, InputIterator last,
                          OutputIterator result);
    template<class InputIterator, class OutputIterator,
              class BinaryOperation>
        OutputIterator
```

```

        inclusive_scan(InputIterator first, InputIterator last,
                       OutputIterator result,
                       BinaryOperation binary_op);
template<class InputIterator, class OutputIterator,
        class T, class BinaryOperation, class T>
OutputIterator
    inclusive_scan(InputIterator first, InputIterator last,
                  OutputIterator result,
                  T init, BinaryOperation binary_op, T init);
}
}
}

```

#### 4.4.2 Reduce

[parallel.alg.reduce]

```

1      template<class InputIterator>
        typename iterator_traits<InputIterator>::value_type
        reduce(InputIterator first, InputIterator last);

```

2 *Effects:* Same as `reduce(first, last, typename iterator_traits<InputIterator>::value_type{}).`

*Returns:* `reduce(first, last, typename iterator_traits<InputIterator>::value_type{})`

*Requires:* `typename iterator_traits<InputIterator>::value_type{} shall be a valid expression.`  
The operator+ function associated with `iterator_traits<InputIterator>::value_type` shall not invalidate iterators or subranges, nor modify elements in the range `[first, last)`.

*Complexity:*  $O(\text{last} - \text{first})$  applications of operator+.

*Notes:* The primary difference between `reduce` and `accumulate` is that the behavior of `reduce` may be non-deterministic for non-associative or non-commutative operator+.

```

3      template<class InputIterator, class T>
        T reduce(InputIterator first, InputIterator last, T init);

```

4 *Effects:* Same as `reduce(first, last, init, plus<>()).`

*Returns:* `reduce(first, last, init, plus<>())`

*Requires:* The operator+ function associated with `T` shall not invalidate iterators or subranges, nor modify elements in the range `[first, last)`.

*Complexity:*  $O(\text{last} - \text{first})$  applications of operator+.

*Notes:* The primary difference between `reduce` and `accumulate` is that the behavior of `reduce` may be non-deterministic for non-associative or non-commutative operator+.

- 5       template<class InputIterator, class T, class BinaryOperation>  
           T reduce(InputIterator first, InputIterator last, T init,  
                   BinaryOperation binary\_op);
- 6   *Returns:* *GENERALIZED\_SUM*(binary\_op, init, \*first, ..., ~~\*(first + last - first - 1)\*(first +~~  
~~(last - first) - 1)).~~
- 7   *Requires:* binary\_op shall not invalidate iterators or subranges, nor modify elements in the  
 range [first,last).
- 8   *Complexity:* O(last - first) applications of binary\_op.
- 9   *Notes:* The primary difference between reduce and accumulate is that the behavior of reduce  
 may be non-deterministic for non-associative or non-commutative ~~operator+~~binary\_op.

#### 4.4.3 Exclusive scan

[\[parallel.alg.exclusive.scan\]](#)

- 1       template<class InputIterator, class OutputIterator,  
                   class T>  
           OutputIterator  
           exclusive\_scan(InputIterator first, InputIterator last,  
                           OutputIterator result,  
                           T init);
- 2   *Effects:* Same as exclusive\_scan(first, last, result, init, plus<>()).

*Returns:* exclusive\_scan(first, last, result, init, plus<>())

*Requires:* The operator+ function associated with iterator\_traits<InputIterator>::value\_type  
 shall not invalidate iterators or subranges, nor modify elements in the ranges [first,last) or  
 [result,result + (last - first)).

*Complexity:* O(last - first) applications of operator+.

*Notes:* The primary difference between exclusive\_scan and inclusive\_scan is that  
 exclusive\_scan excludes the ith input element from the ith sum. If the operator+ function is  
 not mathematically associative, the behavior of exclusive\_scan may be non-deterministic.

- 3       template<class InputIterator, class OutputIterator,  
                  class T, class BinaryOperation>  
          OutputIterator  
          exclusive\_scan(InputIterator first, InputIterator last,  
                          OutputIterator result,  
                          T init, BinaryOperation binary\_op);
- 4   *Effects:* Assigns through each iterator *i* in  $[result, result + (last - first))$  the value of  
 $GENERALIZED\_NONCOMMUTATIVE\_SUM(binary\_op, init, *first, \dots, (*first + i - result - 1) * (first +$   
 $(i - result) - 1))$ .
- 5   *Returns:* The end of the resulting range beginning at *result*.
- 6   *Requires:* *binary\_op* shall not invalidate iterators or subranges, nor modify elements in the  
 ranges  $[first, last)$  or  $[result, result + (last - first))$ .
- 7   *Complexity:*  $O(last - first)$  applications of *binary\_op*.
- 8   *Notes:* The **primary** difference between *exclusive\_scan* and *inclusive\_scan* is that  
*exclusive\_scan* excludes the *i*th input element from the *i*th sum. If *binary\_op* is not  
 mathematically associative, the behavior of *exclusive\_scan* may be non-deterministic.

#### 4.4.4 Inclusive scan

[parallel.alg.inclusive.scan]

- 1       template<class InputIterator, class OutputIterator>  
          OutputIterator  
          inclusive\_scan(InputIterator first, InputIterator last,  
                          OutputIterator result);
- 2   *Effects:* Same as *inclusive\_scan(first, last, result, plus<>())*.
- Returns:* *inclusive\_scan(first, last, result, plus<>())*
- Requires:* The *operator+* function associated with *iterator\_traits<InputIterator>::value\_type*  
 shall not invalidate iterators or subranges, nor modify elements in the ranges  $[first, last)$  or  
 $[result, result + (last - first))$ .
- Complexity:*  $O(last - first)$  applications of *operator+*.
- Notes:* The primary difference between *exclusive\_scan* and *inclusive\_scan* is that  
*exclusive\_scan* excludes the *i*th input element from the *i*th sum. If the *operator+* function is  
 not mathematically associative, the behavior of *inclusive\_scan* may be non-deterministic.

```

3      template<class InputIterator, class OutputIterator,
          class BinaryOperation>
          OutputIterator
          inclusive_scan(InputIterator first, InputIterator last,
                        OutputIterator result,
                        BinaryOperation binary_op);

      template<class InputIterator, class OutputIterator,
          class T, class BinaryOperation, class T>
          OutputIterator
          inclusive_scan(InputIterator first, InputIterator last,
                        OutputIterator result,
                        T init, BinaryOperation binary_op, T init);

```

- <sup>4</sup> *Effects:* Assigns through each iterator *i* in  $[result, result + (last - first))$  the value of *GENERALIZED\_NONCOMMUTATIVE\_SUM*(*binary\_op*, *\*first*, ...,  ~~$(*first + i - result) * (first + (i - result))$~~ ) or *GENERALIZED\_NONCOMMUTATIVE\_SUM*(*binary\_op*, *init*, *\*first*, ...,  ~~$(*first + i - result) * (first + (i - result))$~~ ) if *init* is provided.
- <sup>5</sup> *Returns:* The end of the resulting range beginning at *result*.
- <sup>6</sup> *Requires:* *binary\_op* shall not invalidate iterators or subranges, nor modify elements in the ranges  $[first, last)$  or  $[result, result + (last - first))$ .
- <sup>7</sup> *Complexity:*  $O(last - first)$  applications of *binary\_op*.
- <sup>8</sup> *Notes:* The **primary** difference between *exclusive\_scan* and *inclusive\_scan* is that *inclusive\_scan* includes the *i*th input element in the *i*th sum. If *binary\_op* is not mathematically associative, the behavior of *inclusive\_scan* may be non-deterministic.